



Published in Towards Data Science



Simeon Kostadinov

Follow

Aug 8, 2019 · 8 min read · Listen



Save



# Understanding Backpropagation Algorithm

Learn the nuts and bolts of a neural network's most important ingredient



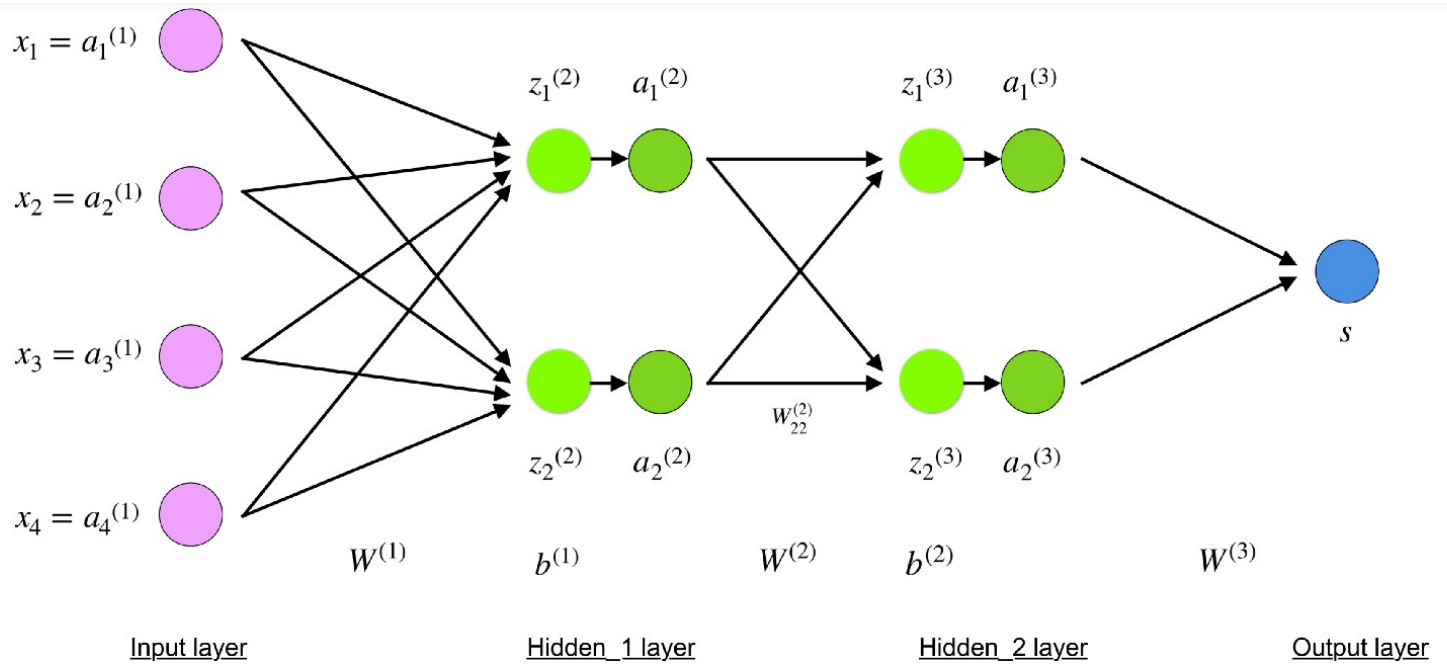
"A man is running on a highway" — photo by [Andrea Leopardi](#) on [Unsplash](#)

**Backpropagation algorithm** is probably the most fundamental building block in a neural network. It was first introduced in 1960s and almost 30 years later (1989) popularized by Rumelhart, Hinton and Williams in a paper called "[\*Learning representations by back-propagating errors\*](#)".

**The algorithm is used to effectively train a neural network through a method called chain rule.** In simple terms, after each forward pass through a network, backpropagation performs a backward pass while adjusting the model's parameters (weights and biases).

In this article, I would like to go over the mathematical process of training and optimizing a simple 4-layer neural network. *I believe this would help the reader understand how backpropagation works as well as realize its importance.*





### Input layer

The neurons, colored in **purple**, represent the input data. These can be as simple as scalars or more complex like vectors or multidimensional matrices.

$$x_i = a_i^{(1)}, i \in 1, 2, 3, 4$$

Equation for input  $x_i$

The first set of activations ( $a$ ) are equal to the input values. NB: “activation” is the neuron’s value after applying an activation function. See below.

### Hidden layers

The final values at the hidden neurons, colored in **green**, are computed using  $z^l$  — weighted inputs in layer  $l$ , and  $a^l$  — activations in layer  $l$ . For layer 2 and 3 the equations are:

- $l = 2$

$$z^{(2)} = W^{(1)}x + b^{(1)}$$

$$a^{(2)} = f(z^{(2)})$$

Equations for  $z^2$  and  $a^2$

- $l = 3$

$$z^{(3)} = W^{(2)}a^{(2)} + b^{(2)}$$

$$a^{(3)} = f(z^{(3)})$$





Activations  $a^L$  and  $a^J$  are computed using an activation function  $f$ . Typically, this **function  $f$  is non-linear** (e.g. [sigmoid](#), [ReLU](#), [tanh](#)) and allows the network to learn complex patterns in data. We won't go over the details of how activation functions work, but, if interested, I strongly recommend reading [this great article](#).

Looking carefully, you can see that all of  $x$ ,  $z^2$ ,  $a^2$ ,  $z^3$ ,  $a^3$ ,  $W^1$ ,  $W^2$ ,  $b^1$  and  $b^2$  are missing their subscripts presented in the 4-layer network illustration above. **The reason is that we have combined all parameter values in matrices, grouped by layers.** This is the standard way of working with neural networks and one should be comfortable with the calculations. However, I will go over the equations to clear out any confusion.

Let's pick layer 2 and its parameters as an example. The same operations can be applied to any layer in the network.

- $W^1$  is a weight matrix of shape  $(n, m)$  where  $n$  is the number of output neurons (neurons in the next layer) and  $m$  is the number of input neurons (neurons in the previous layer). For us,  $n = 2$  and  $m = 4$ .

$$W^{(1)} = \begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} & W_{14}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} & W_{24}^{(1)} \end{bmatrix}$$

Equation for  $W^1$

**NB: The first number in any weight's subscript matches the index of the neuron in the next layer** (in our case this is the *Hidden\_2 layer*) **and the second number matches the index of the neuron in previous layer** (in our case this is the *Input layer*).

- $x$  is the input vector of shape  $(m, 1)$  where  $m$  is the number of input neurons. For us,  $m = 4$ .

$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}$$

Equation for  $x$

- $b^1$  is a bias vector of shape  $(n, 1)$  where  $n$  is the number of neurons in the current layer. For us,  $n = 2$ .


$$b^{(1)} = \begin{bmatrix} b_1^{(1)} \\ b_2^{(1)} \end{bmatrix}$$

Equation for  $b^1$

Following the equation for  $z^2$ , we can use the above definitions of  $W^1$ ,  $x$  and  $b^1$  to derive "Equation for  $z^2$ ":

$$z^{(2)} = \begin{bmatrix} W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + W_{13}^{(1)}x_3 + W_{14}^{(1)}x_4 \\ W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + W_{24}^{(1)}x_4 \end{bmatrix} + \begin{bmatrix} b_1^{(1)} \\ b_2^{(1)} \end{bmatrix}$$





You will see that  $z^2$  can be expressed using  $(z_1)^2$  and  $(z_2)^2$  where  $(z_1)^2$  and  $(z_2)^2$  are the sums of the multiplication between every input  $x_i$  with the corresponding weight  $(W_{ij})^1$ .

This leads to the same “Equation for  $z^2$ ” and proves that the matrix representations for  $z^2$ ,  $a^2$ ,  $z^3$  and  $a^3$  are correct.

### Output layer

The final part of a neural network is the output layer which produces the predicated value. In our simple example, it is presented as a single neuron, colored in **blue** and evaluated as follows:

$$s = W^{(3)}a^{(3)}$$

Equation for output s

Again, we are using the matrix representation to simplify the equation. One can use the above techniques to understand the underlying logic. **Please leave any comments below if you find yourself lost in the equations — I would love to help!**

### Forward propagation and evaluation

The equations above form network’s forward propagation. Here is a short overview:





$$z^{(2)} = W^{(1)}x + b^{(1)} \quad \text{neuron value at Hidden}_1 \text{ layer}$$

$$a^{(2)} = f(z^{(2)}) \quad \text{activation value at Hidden}_1 \text{ layer}$$

$$z^{(3)} = W^{(2)}a^{(2)} + b^{(2)} \quad \text{neuron value at Hidden}_2 \text{ layer}$$

$$a^{(3)} = f(z^{(3)}) \quad \text{activation value at Hidden}_2 \text{ layer}$$

$$s = W^{(3)}a^{(3)} \quad \text{Output layer}$$

Overview of forward propagation equations colored by layer

The final step in a forward pass is to evaluate the **predicted output**  $s$  against an **expected output**  $y$ .

The output  $y$  is part of the training dataset  $(x, y)$  where  $x$  is the input (as we saw in the previous section).

Evaluation between  $s$  and  $y$  happens through a **cost function**. This can be as simple as MSE (mean squared error) or more complex like cross-entropy.

We name this cost function  $C$  and denote it as follows:

$$C = cost(s, y)$$

Equation for cost function  $C$

where  $cost$  can be equal to MSE, cross-entropy or any other cost function.

Based on  $C$ 's value, the model “knows” how much to adjust its parameters in order to get closer to the expected output  $y$ . This happens using the backpropagation algorithm.

### Backpropagation and computing gradients

According to the paper from 1989, backpropagation:

*repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector.*

and

*the ability to create useful new features distinguishes back-propagation from earlier, simpler methods...*

In other words, **backpropagation aims to minimize the cost function by adjusting network's weights and biases**. The level of adjustment is determined by the gradients of the cost function with respect to those parameters.

One question may arise — **why computing gradients?**





- Gradient of a function  $C(x_1, x_2, \dots, x_m)$  in point  $x$  is a vector of the partial derivatives of  $C$  in  $x$ .

$$\frac{\partial C}{\partial x} = \left[ \frac{\partial C}{\partial x_1}, \frac{\partial C}{\partial x_2}, \dots, \frac{\partial C}{\partial x_m} \right]$$

Equation for derivative of  $C$  in  $x$

- The derivative of a function  $C$  measures the sensitivity to change of the function value (output value) with respect to a change in its argument  $x$  (input value). In other words, the derivative tells us the direction  $C$  is going.
- The gradient shows how much the parameter  $x$  needs to change (in positive or negative direction) to minimize  $C$ .

Compute those gradients happens using a technique called chain rule.

For a single weight  $(w_{jk})^l$ , the gradient is:

$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l} \quad \text{chain rule}$$

$$z_j^l = \sum_{k=1}^m w_{jk}^l a_k^{l-1} + b_j^l \quad \text{by definition}$$

$m$  – number of neurons in  $l-1$  layer

$$\frac{\partial z_j^l}{\partial w_{jk}^l} = a_k^{l-1} \quad \text{by differentiation (calculating derivative)}$$

$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} a_k^{l-1} \quad \text{final value}$$

Equations for derivative of  $C$  in a single weight  $(w_{jk})^l$

Similar set of equations can be applied to  $(b_j)^l$ :







$$\frac{\partial b_j^l}{\partial z_j^l} = \frac{\partial z_j^l}{\partial b_j^l} \frac{\partial b_j^l}{\partial z_j^l} \quad \text{chain rule}$$

$$\frac{\partial z_j^l}{\partial b_j^l} = 1 \quad \text{by differentiation (calculating derivative)}$$

$$\frac{\partial C}{\partial b_j^l} = \frac{\partial C}{\partial z_j^l} 1 \quad \text{final value}$$

Equations for derivative of C in a single bias  $(b_j)^l$

The common part in both equations is often called “local gradient” and is expressed as follows:

$$\delta_j^l = \frac{\partial C}{\partial z_j^l} \quad \text{local gradient}$$

Equation for local gradient

The “local gradient” can easily be determined using the chain rule. I won’t go over the process now but if you have any questions, please comment below.

The gradients allow us to optimize the model’s parameters:

*while (termination condition not met)*

$$w := w - \epsilon \frac{\partial C}{\partial w}$$

$$b := b - \epsilon \frac{\partial C}{\partial b}$$

*end*


Algorithm for optimizing weights and biases (also called “Gradient descent”)

- Initial values of  $w$  and  $b$  are randomly chosen.
- Epsilon ( $\epsilon$ ) is the learning rate. It determines the gradient’s influence.
- $w$  and  $b$  are matrix representations of the weights and biases. Derivative of  $C$  in  $w$  or  $b$  can be calculated using partial derivatives of  $C$  in the individual weights or biases.
- Termination condition is met once the cost function is minimized.





Let's zoom in on the bottom part of the above neural network:



Visual representation of backpropagation in a neural network

Weight  $(w_{22})^2$  connects  $(a_2)^2$  and  $(z_2)^2$ , so computing the gradient requires applying the chain rule through  $(z_2)^3$  and  $(a_2)^3$ :

$$\frac{\partial C}{\partial w_{22}^{(2)}} = \frac{\partial C}{\partial z_2^{(3)}} \cdot \frac{\partial z_2^{(3)}}{\partial w_{22}^{(2)}} = \frac{\partial C}{\partial a_2^{(3)}} \cdot \frac{\partial a_2^{(3)}}{\partial z_2^{(3)}} \cdot a_2^{(2)} = \frac{\partial C}{\partial a_2^{(3)}} \cdot f'(z_2^{(3)}) \cdot a_2^{(2)}$$

Equation for derivative of C in  $(w_{22})^2$

Calculating the final value of derivative of C in  $(a_2)^3$  requires knowledge of the function C. Since C is dependent on  $(a_2)^3$ , calculating the derivative should be fairly straightforward.

I hope this example manages to throw some light on the mathematics behind computing gradients. To further enhance your skills, I strongly recommend watching [Stanford's NLP series where Richard Socher gives 4 great explanations of backpropagation](#).

### Final remarks

In this article, I went through a detailed explanation of how backpropagation works under the hood using mathematical techniques like computing gradients, chain rule etc. *Knowing the nuts and bolts of this algorithm will fortify your neural networks knowledge and make you feel comfortable to take on more complex models. Enjoy your deep learning journey!*

**Thank you for the reading. Hope you enjoyed the article 🥰 and I wish you a great day!**

### Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to







Get unlimited access Open in app





Get unlimited access Open in app

