

[Open in app](#)[Get started](#)

Published in MLearning.ai



Andrea D'Agostino

[Follow](#)Nov 24, 2021 · 7 min read · [Listen](#)

# Text Clustering with TF-IDF in Python

Explanation of a simple pipeline for text clustering. Full example and code



Photo by [Andrew Wulf](#) on [Unsplash](#)

TF-IDF is a well known and documented **vectorization technique** in data science. Vectorization is the act of converting data into a numerical format in such a way that a



[Open in app](#)[Get started](#)

## Methodology

We will use a dataset provided by Sklearn to have a replicable corpus. After that, we will use the KMeans algorithm to group the vectors generated by the TF-IDF. We will then use Principal Component Analysis to visualize our groups and bring out common or unusual characteristics of the texts present in our corpus.

Here is what we'll do

- import the dataset
- apply **preprocessing** to our corpus to remove words and symbols which, when converted into numerical format, do not add value to our model
- use **TF-IDF** as a vectorization algorithm
- apply **KMeans** to group our data
- apply **PCA** to reduce the dimensionality of our vectors to 2 for visualization purposes
- **interpret** the data

## The Analysis

### Our Dataset

For this example we will use Scikit-Learn's API, *sklearn.datasets*, which allows us to access a famous dataset for linguistic analysis, the **20newsgroups dataset**. A newsgroup is an online user discussion group, such as a forum. Sklearn allows us to access different categories of content. We will use texts that have to do with technology, religion and sport.

```
1 # import the dataset from sklearn
2 from sklearn.datasets import fetch_20newsgroups
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 from sklearn.cluster import KMeans
```



[Open in app](#)[Get started](#)

```
10
11 # string manipulation libs
12 import re
13 import string
14 import nltk
15 from nltk.corpus import stopwords
16
17 # viz libs
18 import matplotlib.pyplot as plt
19 import seaborn as sns
20
21
22 categories = [
23     'comp.graphics',
24     'comp.os.ms-windows.misc',
25     'rec.sport.baseball',
26     'rec.sport.hockey',
27     'alt.atheism',
28     'soc.religion.christian',
29 ]
30 dataset = fetch_20newsgroups(subset='train', categories=categories, shuffle=True, remove=)
```

clustering\_eng\_1.py hosted with ❤ by GitHub

[view raw](#)

If we access the first element with `dataset['data'][0]` we see

*They tried their best not to show it, believe me. I'm surprised they couldn't find a sprint car race (mini cars through pigpens, indeed!) On short notice.*

George

It is possible that your data is different due to `shuffle=True`, which randomizes the order of the elements of the dataset. The number of items in our dataset is 3451.

Let's create a Pandas dataframe from our dataset

```
1 df = pd.DataFrame(dataset.data, columns=["corpus"])
```




[Open in app](#)
[Get started](#)
**corpus**

0 \nThey tried their best not to show it, believ...

1 \nStankiewicz? I doubt it.\n\nKoufax was one ...

2 \n[deletia- and so on]\n\nI seem to have been ...

3 Excuse the sheer newbiness of this post, but ...

4 =====...

...

3446 \n Or, with no dictionary available, they cou...

3447 \n\nSorry to disappoint you but the Red Wings ...

3448 \n: Can anyone tell me where to find a MPEG vi...

3449 \n

3450 \nHey Valentine, I don't see Boston with any w...

3451 rows x 1 columns

Our corpus before the cleaning phase

Note that there are `\n`, `===` and other symbols that should be removed to properly train our model.

## Preprocessing

Along with the symbols mentioned, we also want remove stopwords . These are a series of words that add no information to our model. An example of a stopwords in English are articles, conjunctions, and so on.

We will use the NLTK library to import the stopwords and display them



[Open in app](#)[Get started](#)

```
5 stopwords.words("english")[:10] # <-- import the english stopwords
```

clustering\_eng\_4.py hosted with ❤ by GitHub

[view raw](#)

>>> ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're"] and so on.

Now let's create a *preprocess\_text* function that takes some text as input and returns a clean version of it.

```
1 def preprocess_text(text: str, remove_stopwords: bool) -> str:
2     """This utility function sanitizes a string by:
3     - removing links
4     - removing special characters
5     - removing numbers
6     - removing stopwords
7     - transforming in lowercase
8     - removing excessive whitespaces
9     Args:
10         text (str): the input text you want to clean
11         remove_stopwords (bool): whether or not to remove stopwords
12     Returns:
13         str: the cleaned text
14     """
15
16     # remove links
17     text = re.sub(r"http\S+", "", text)
18     # remove special chars and numbers
19     text = re.sub("[^A-Za-z]+", " ", text)
20     # remove stopwords
21     if remove_stopwords:
22         # 1. tokenize
23         tokens = nltk.word_tokenize(text)
24         # 2. check if stopword
25         tokens = [w for w in tokens if not w.lower() in stopwords.words("english")]
26         # 3. join back together
27         text = " ".join(tokens)
28     # return text in lower case and stripped of whitespaces
29     text = text.lower().strip()
```




[Open in app](#)
[Get started](#)

Here is the same previous document, preprocessed

*tried best show believe im surprised couldnt find sprint car race mini cars pigpens  
indeed short notice george*

Let's apply the function to the entire dataframe

```
1 df['cleaned'] = df['corpus'].apply(lambda x: preprocess_text(x, remove_stopwords=True))
```

clustering\_5.py hosted with ❤ by GitHub

[view raw](#)

	corpus	cleaned
0	\nThey tried their best not to show it, believ...	they tried their best not to show it believe m...
1	\nStankiewicz? I doubt it.\n\nKoufax was one ...	stankiewicz i doubt it koufax was one of two j...
2	\n[deletia- and so on]\n\nI seem to have been ...	deletia and so on i seem to have been rather u...
3	Excuse the sheer newbiness of this post, but ...	excuse the sheer newbiness of this post but i...
4	=====...	
...	...	...
3446	\n Or, with no dictionary available, they cou...	or with no dictionary available they could gai...
3447	\n\nSorry to disappoint you but the Red Wings ...	sorry to disappoint you but the red wings earn...
3448	\n: Can anyone tell me where to find a MPEG vi...	can anyone tell me where to find a mpeg viewer...
3449	\n	
3450	\nHey Valentine, I don't see Boston with any w...	hey valentine i dont see boston with any world...


Cleaned corpus series added to our dataframe

Now we are ready to perform vectorization.

## TF-IDF Vectorization

The TF-IDF converts our corpus into a numerical format by bringing out specific terms, weighing very rare or very common terms differently in order to assign them a low score.



[Open in app](#)[Get started](#)

TF-IDF: the term in green gets a high score because it is more specific

With this vectorization technique we are able to group our documents considering the most important terms that constitute them. To read more about how TF-IDF works, read [here](#).

With Sklearn, applying TF-IDF is trivial

```
1 # initialize the vectorizer
2 vectorizer = TfidfVectorizer(sublinear_tf=True, min_df=5, max_df=0.95)
3 # fit_transform applies TF-IDF to clean texts - we save the array of vectors in X
4 X = vectorizer.fit_transform(df['cleaned'])
```

clustering\_eng\_vec.py hosted with ❤ by GitHub

[view raw](#)

X is the array of vectors that will be used to train the KMeans model. The default






[Open in app](#)
[Get started](#)

```
vector_a = [1.204, 0, 0, 0, 0, 0, 0, ..., 0]
```

The vector is made up of a single value not equal to 0. In Sklearn, a sparse matrix is nothing more than a matrix that indexes the position of values and 0s instead of storing it like any other matrix. This is a mechanism for saving RAM and computing power. The convenience is that the sparse matrix is accepted by most machine learning algorithms, as well as KMeans. In fact, the latter will use the data present in the sparse matrix to find groups and patterns.

If we use `X.toarray()` we actually see the complete matrix, not sparse.

```
array([[0., 0., 0., ..., 0., 0., ],
       [0., 0., 0., ..., 0., 0., ],
       [0., 0., 0., ..., 0., 0., ],
       [0.10747419, 0., 0., ..., 0., 0., ],
       [0., 0., 0., ..., 0., 0., ],
       ...,
       [0., 0., 0., ..., 0., 0., ],
       [0., 0., 0., ..., 0., 0., ],
       [0., 0., 0., ..., 0., 0., ],
       [0., 0., 0., ..., 0., 0., ],
       [0., 0., 0., ..., 0., 0., ]])
```


Vector representation of the sparse matrix

## Implementation of KMeans

KMeans is one of the most known and used unsupervised algorithms in data science and is used to group a set of data into a defined number of groups. The idea behind it is very simple: the algorithm initializes random positions (called centroids, the red, blue and green points in the screenshot below) in the vector plane and assigns the point to the nearest centroid.






[Open in app](#)[Get started](#)

First step of KMeans — initialization of the centroids

Screenshot taken from [https://www.youtube.com/watch?v=R2e3Ls9H\\_fc](https://www.youtube.com/watch?v=R2e3Ls9H_fc)

The algorithm calculates the average position (or, if it helps the interpretation, the “center of gravity”) of the points and moves the respective centroid to that position and updates the group to which each point belongs. The algorithm converges when all points are at the minimum distance from their respective centroid.



[Open in app](#)[Get started](#)

## Convergence of the KMeans — cluster discovery

Screenshot taken from [https://www.youtube.com/watch?v=R2e3Ls9H\\_fc](https://www.youtube.com/watch?v=R2e3Ls9H_fc)

Let's continue with our project by going to use Sklearn again

```
1 from sklearn.cluster import KMeans
2
3 # initialize kmeans with 3 centroids
4 kmeans = KMeans(n_clusters=3, random_state=42)
5 # fit the model
6 kmeans.fit(X)
7 # store cluster labels in a variable
8 clusters = kmeans.labels_
```

clustering\_eng\_kmeans.py hosted with ❤ by GitHub

[view raw](#)

With this code snippet we trained the KMeans with the vectors returned by the TF-IDF and we assigned the groups to the *clusters* variable.

```
1 [c for c in clusters][:10]
[0, 0, 2, 1, 0, 0, 0, 1, 1, 1]
```

Now we can proceed to the visualization of our groups and evaluate their segmentation and / or presence of anomalies.

## Dimensional Reduction and Visualization

We have our X from the TF-IDF and we have a KMeans model and related clusters. Now we want to put these two pieces together to visualize the relation between text and group.

As we know, a graph is usually presented in 2 dimensions and rarely in 3. Surely we cannot visualize more. If we look at the dimensionality of X with *X.shape* we see that it is



[Open in app](#)[Get started](#)

Fortunately for us, there is a technique called **PCA (Principal Component Analysis)** which reduces the dimensionality of a data set to an arbitrary number while preserving most of the information contained in it. As with the TF-IDF, I won't go into detail about how it works, and you can read more about how it works [here](#).

For the purpose of this article, it is enough for us to know that PCA tends to preserve the dimensions that best summarize the total variability of our data, by removing dimensions that contribute little to the latter.

Our X will go from 7390 in size to 2. *Sklearn.decomposition.PCA* is what we need.

```
1 from sklearn.decomposition import PCA
2
3 # initialize PCA with 2 components
4 pca = PCA(n_components=2, random_state=42)
5 # pass our X to the pca and store the reduced vectors into pca_vecs
6 pca_vecs = pca.fit_transform(X.toarray())
7 # save our two dimensions into x0 and x1
8 x0 = pca_vecs[:, 0]
9 x1 = pca_vecs[:, 1]
```

clustering\_eng\_pca.py hosted with ❤ by GitHub

[view raw](#)

1 x0

```
array([-0.00152024, -0.03644996, -0.06548033, ..., 0.18883194,
       0.03557314, -0.05786917])
```

1 x1

```
array([-0.00430002, -0.03914495, 0.08785332, ..., 0.05718469,
       -0.03828756, -0.10444227])
```

Two two reduced dimensions generated by the PCA algorithm




[Open in app](#)
[Get started](#)

Before creating our chart let's better organize our dataframe by creating columns *cluster*, *x0*, *x1*

```
1 # assign clusters and pca vectors to our dataframe
2 df['cluster'] = clusters
3 df['x0'] = x0
4 df['x1'] = x1
```

clustering\_eng\_pca.py hosted with ❤ by GitHub

[view raw](#)

	corpus	cleaned	cluster	x0	x1
0	\nThey tried their best not to show it, believ...	tried best show believe surprised find sprint ...	0	-0.001520	-0.004300
1	\nStankiewicz? I doubt it.\n\nKoufax was one ...	stankiewicz doubt koufax one two jewish hofs h...	0	-0.036450	-0.039145
2	\n[deletia- and so on]\n\nI seem to have been ...	deletia seem rather unclear asking please show...	2	-0.065480	0.087853
3	Excuse the sheer newbiness of this post, but ...	excuse sheer newbiness post looking decent pa...	1	0.164120	0.062958
4	=====...		0	0.035573	-0.038288
...	...	...	...	...	...
3446	\n Or, with no dictionary available, they cou...	dictionary available could gain first hand kno...	0	-0.010061	-0.010073
3447	\n\nSorry to disappoint you but the Red Wings ...	sorry disappoint red wings earned victory easi...	0	-0.029921	-0.158443
3448	\n: Can anyone tell me where to find a MPEG vi...	anyone tell find mpeg viewer either dos window...	1	0.188832	0.057185
3449	\n		0	0.035573	-0.038288
3450	\nHey Valentine, I don't see Boston with any w...	hey valentine see boston world series rings fi...	0	-0.057869	-0.104442

3451 rows x 5 columns

Dataset ready for visualization after KMeans and PCA application

Let's see which are the most relevant keywords for each centroid so that we can rename each cluster with a better label

```
1 def get_top_keywords(n_terms):
2     """This function returns the keywords for each centroid of the KMeans"""
3     df = pd.DataFrame(X.todense()).groupby(clusters).mean() # groups the TF-IDF vector by c
4     terms = vectorizer.get_feature_names_out() # access tf-idf terms
5     for i,r in df.iterrows():
6         print('\nCluster {}'.format(i))
7         print(', '.join([terms[t] for t in np.argsort(r)[-n_terms:]])) # for each row of the
8
9     get_top_keywords(10)
```



[Open in app](#)[Get started](#)

```
Cluster 0
good, last, games, like, would, year, think, one, team, game

Cluster 1
please, dos, use, know, program, anyone, files, file, thanks, windows

Cluster 2
christians, say, think, bible, believe, jesus, one, would, people, god
```

Keywords per centroid

Good! The KMeans has correctly created 3 distinct groups, one for each category present in the dataset. Cluster 0 refers to sport, cluster 2 to software / tech, cluster 3 to religion. Let's apply the mapping

```
1 # map clusters to appropriate labels
2 cluster_map = {0: "sport", 1: "tech", 2: "religion"}
3 # apply mapping
4 df['cluster'] = df['cluster'].map(cluster_map)
```

clustering\_eng\_mapping.py hosted with ❤️ by GitHub


[view raw](#)

Let's proceed with the Seaborn library to visualize our grouped texts in a very simple way.

```
1 # set image size
2 plt.figure(figsize=(12, 7))
3 # set a title
4 plt.title("TF-IDF + KMeans 20newsgroup clustering", fontdict={"fontsize": 18})
5 # set axes names
6 plt.xlabel("X0", fontdict={"fontsize": 16})
7 plt.ylabel("X1", fontdict={"fontsize": 16})
8 # create scatter plot with seaborn, where hue is the class used to group the data
9 sns.scatterplot(data=df, x='x0', y='x1', hue='cluster', palette="viridis")
10 plt.show()
```

clustering\_eng\_viz.py hosted with ❤️ by GitHub

[view raw](#)

[Open in app](#)[Get started](#)

Text data clustering using TF-IDF and KMeans. Each point is a vectorized text belonging to a defined category

As we can see, the clustering activity worked well: the algorithm found three distinct groups; as they can be found in our dataset. Imagine the power of this approach in finding clusters in unlabeled data :)

## Interpretation

The interpretation is quite simple: there are no particular anomalies, except for the fact that there are texts belonging to the technology category that overlap slightly with those belonging to sport, between the dark blue and bright green border. This is due to the presence of common terms among some of these texts which, when vectorized, obtain equal values for some dimensions.

As a follow-up it would be interesting to investigate how these texts are written and understand if the hypothesized motivation is well founded or not.



[Open in app](#)[Get started](#)

```
2 from sklearn.datasets import fetch_20newsgroups
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 from sklearn.cluster import KMeans
5 from sklearn.decomposition import PCA
6
7 # import other required libs
8 import pandas as pd
9 import numpy as np
10
11 # string manipulation libs
12 import re
13 import string
14 import nltk
15 nltk.download('punkt')
16 nltk.download('stopwords')
17 from nltk.corpus import stopwords
18
19 # viz libs
20 import matplotlib.pyplot as plt
21 import seaborn as sns
22
23 categories = [
24     'comp.graphics',
25     'comp.os.ms-windows.misc',
26     'rec.sport.baseball',
27     'rec.sport.hockey',
28     'alt.atheism',
29     'soc.religion.christian',
30 ]
31 dataset = fetch_20newsgroups(subset='train', categories=categories, shuffle=True, remove=(
32
33 df = pd.DataFrame(dataset.data, columns=["corpus"])
34 df['cleaned'] = df['corpus'].apply(lambda x: preprocess_text(x, remove_stopwords=True))
35
36 # initialize vectorizer
37 vectorizer = TfidfVectorizer(sublinear_tf=True, min_df=5, max_df=0.95)
38 # fit_transform applies TF-IDF to clean texts - we save the array of vectors in X
39 X = vectorizer.fit_transform(df['cleaned'])
40
41 # initialize KMeans with 3 clusters
```





[Open in app](#)[Get started](#)

```
47  pca = PCA(n_components=2, random_state=42)
48  # pass X to the pca
49  pca_vecs = pca.fit_transform(X.toarray())
50  # save the two dimensions in x0 and x1
51  x0 = pca_vecs[:, 0]
52  x1 = pca_vecs[:, 1]
53
54  # assign clusters and PCA vectors to columns in the original dataframe
55  df['cluster'] = clusters
56  df['x0'] = x0
57  df['x1'] = x1
58
59  cluster_map = {0: "sport", 1: "technology", 2: "religion"} # mapping found through get_top
60  df['cluster'] = df['cluster'].map(cluster_map)
61
62  # set image size
63  plt.figure(figsize=(12, 7))
64  # set title
65  plt.title("Raggruppamento TF-IDF + KMeans 20newsgroup", fontdict={"fontsize": 18})
66  # set axes names
67  plt.xlabel("X0", fontdict={"fontsize": 16})
68  plt.ylabel("X1", fontdict={"fontsize": 16})
69  # create scatter plot with seaborn, where hue is the class used to group the data
70  sns.scatterplot(data=df, x='x0', y='x1', hue='cluster', palette="viridis")
71  plt.show()
```

clustering\_eng\_full.py hosted with ❤ by GitHub

[view raw](#)

```
1  def preprocess_text(text: str, remove_stopwords: bool) -> str:
2      """This function cleans the input text by
3          - removing links
4          - removing special chars
5          - removing numbers
6          - removing stopwords
7          - transforming in lower case
8          - removing excessive whitespaces
9      Arguments:
10         text (str): text to clean
11         remove_stopwords (bool): remove stopwords or not
12     Returns:
```



[Open in app](#)[Get started](#)

```
18     text = re.sub("[^A-Za-z]+", " ", text)
19     # remove stopwords
20     if remove_stopwords:
21         # 1. creates tokens
22         tokens = nltk.word_tokenize(text)
23         # 2. checks if token is a stopwords and removes it
24         tokens = [w for w in tokens if not w.lower() in stopwords.words("english")]
25         # 3. joins all tokens again
26         text = " ".join(tokens)
27     # returns cleaned text
28     text = text.lower().strip()
29     return text
30
31 def get_top_keywords(n_terms):
32     """This function returns the keywords for each centroid of the KMeans"""
33     df = pd.DataFrame(X.todense()).groupby(clusters).mean() # groups tf idf vector per cluster
34     terms = vectorizer.get_feature_names_out() # access to tf idf terms
35     for i,r in df.iterrows():
36         print('\nCluster {}'.format(i))
37         print(', '.join([terms[t] for t in np.argsort(r)[-n_terms:]))) # for each row of the matrix
38
```

clustering\_eng\_functions.py hosted with ❤ by GitHub

[view raw](#)

## MLearning.ai Submission Suggestions

How to become a writer on MLearning.ai

[medium.com](https://medium.com)



[Open in app](#)[Get started](#)

Be sure to SUBSCRIBE here  to never miss another article on Machine Learning & AI Art [Take a look.](#)

Your email

---

Get this newsletter

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

